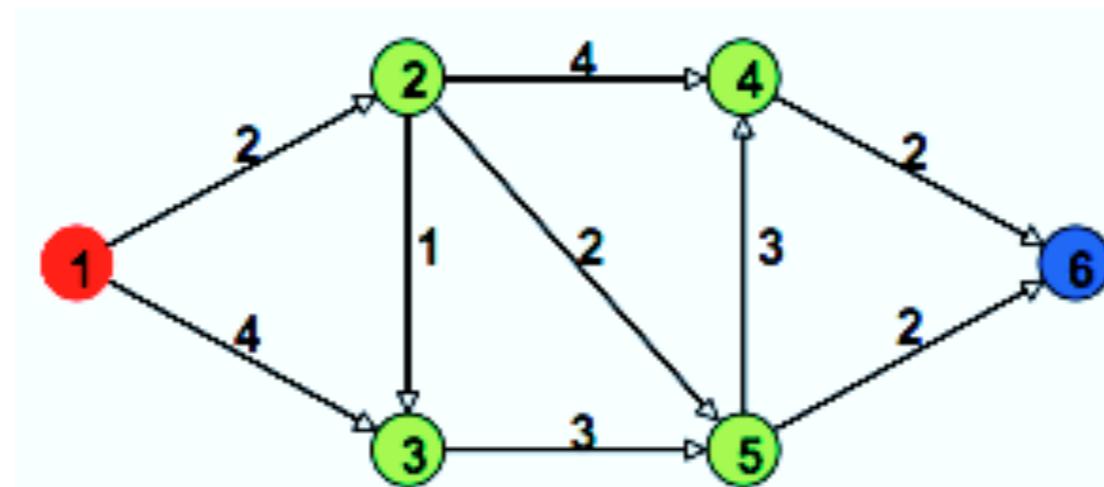
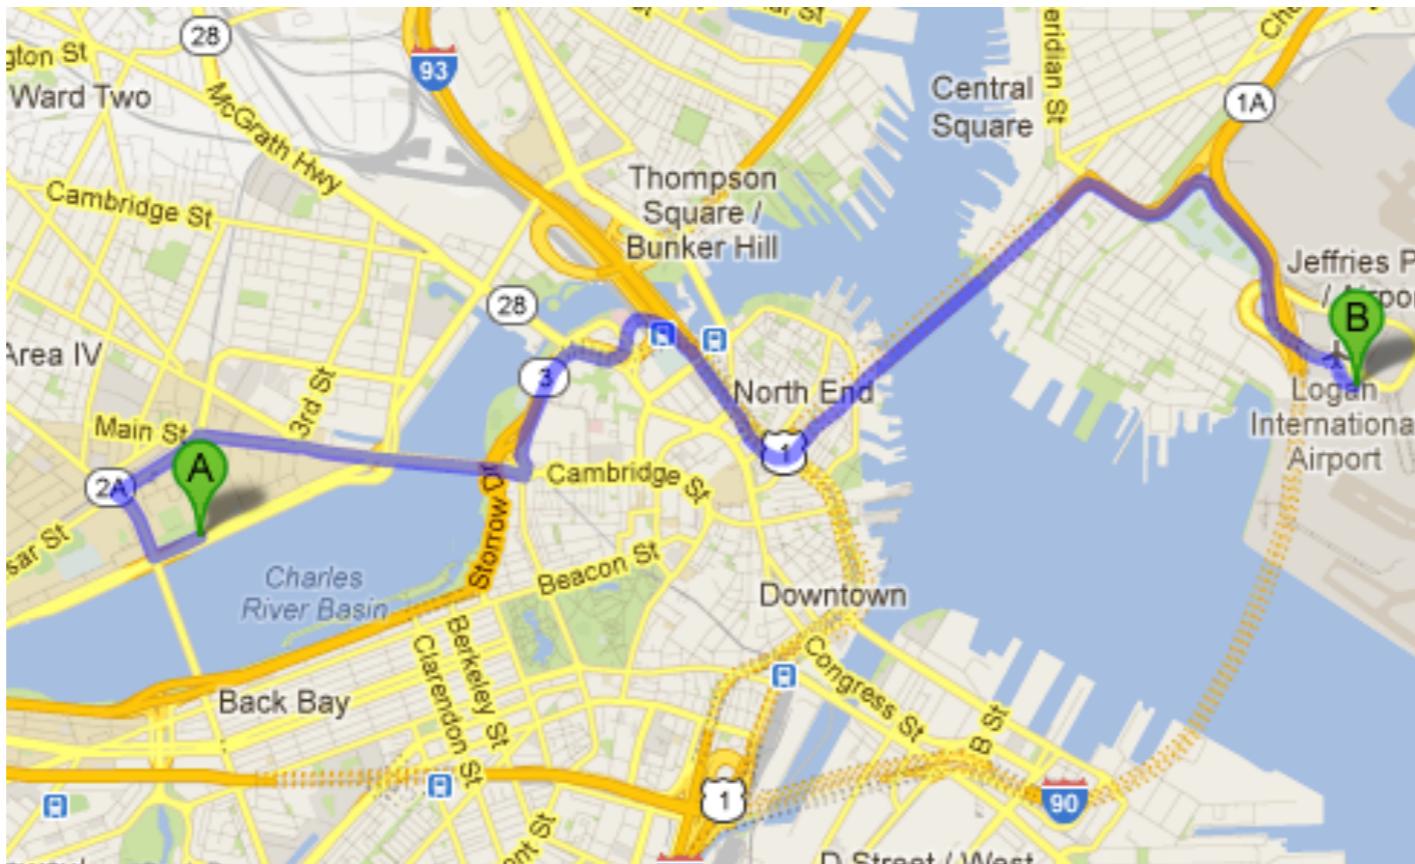
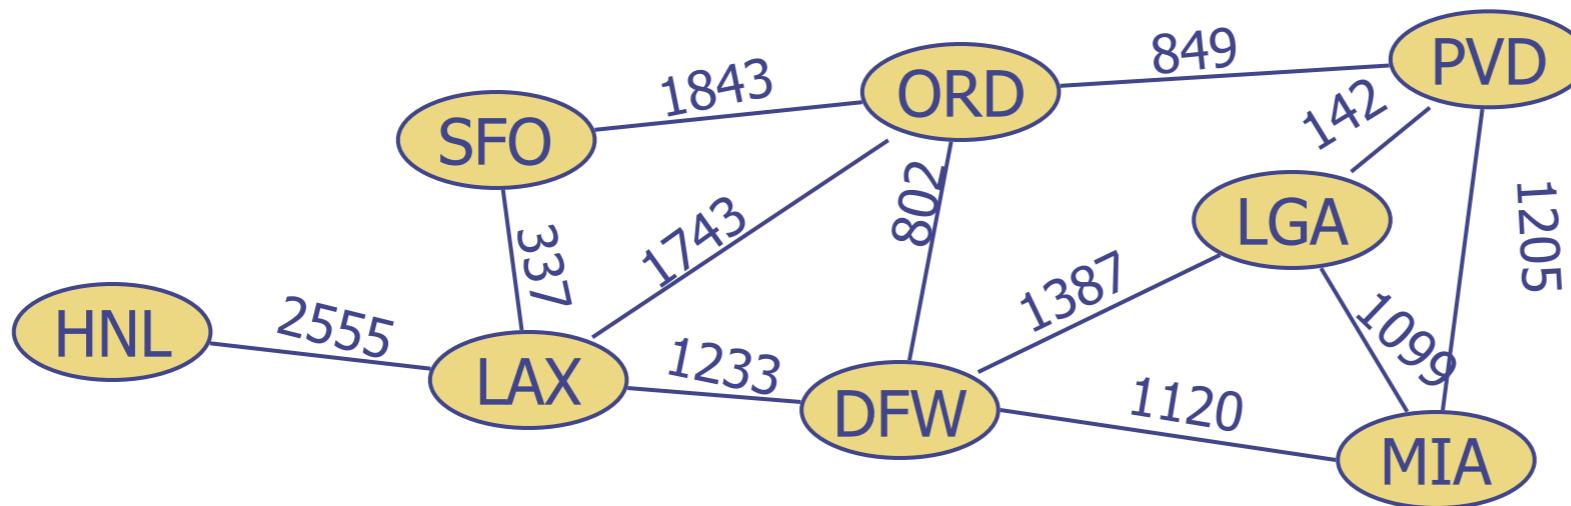


How to get to Logan airport?



Weighted Graphs

- In a weighted graph, each edge has an associated numerical value, called the weight of the edge
- Edge weights may represent, distances, costs, etc.
- Example:
 - In a flight route graph, the weight of an edge represents the distance in miles between the endpoint airports



Shortest Path Properties

Property 1:

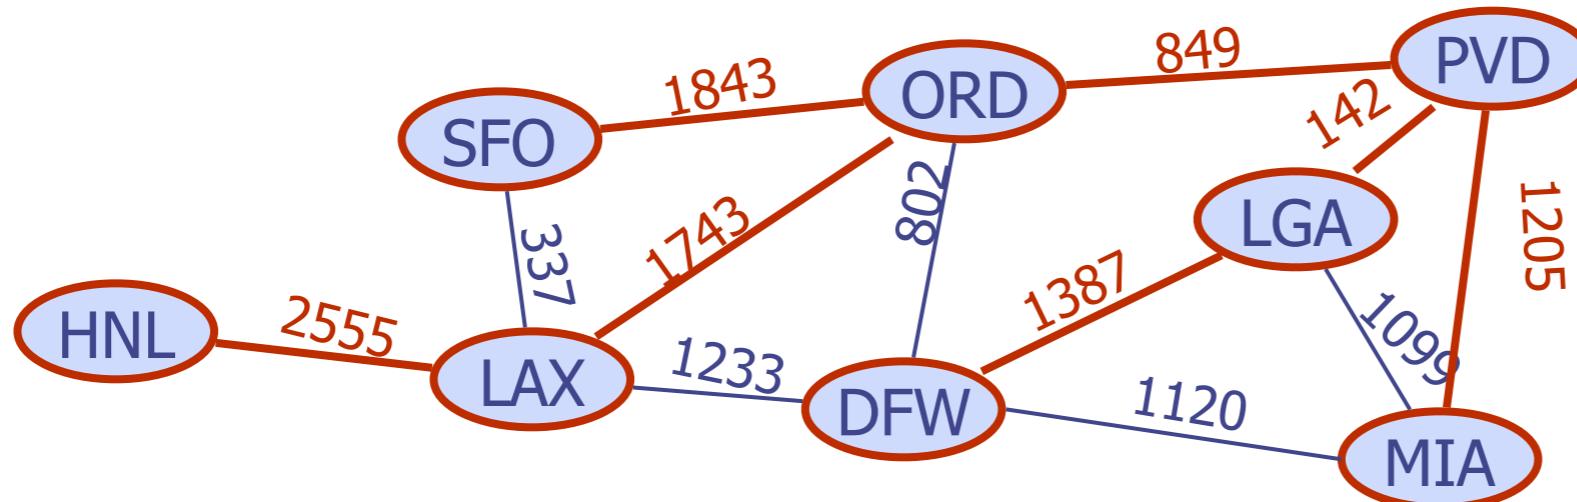
A subpath of a shortest path is itself a shortest path

Property 2:

There is a **tree of shortest paths** from a start vertex to all the other vertices

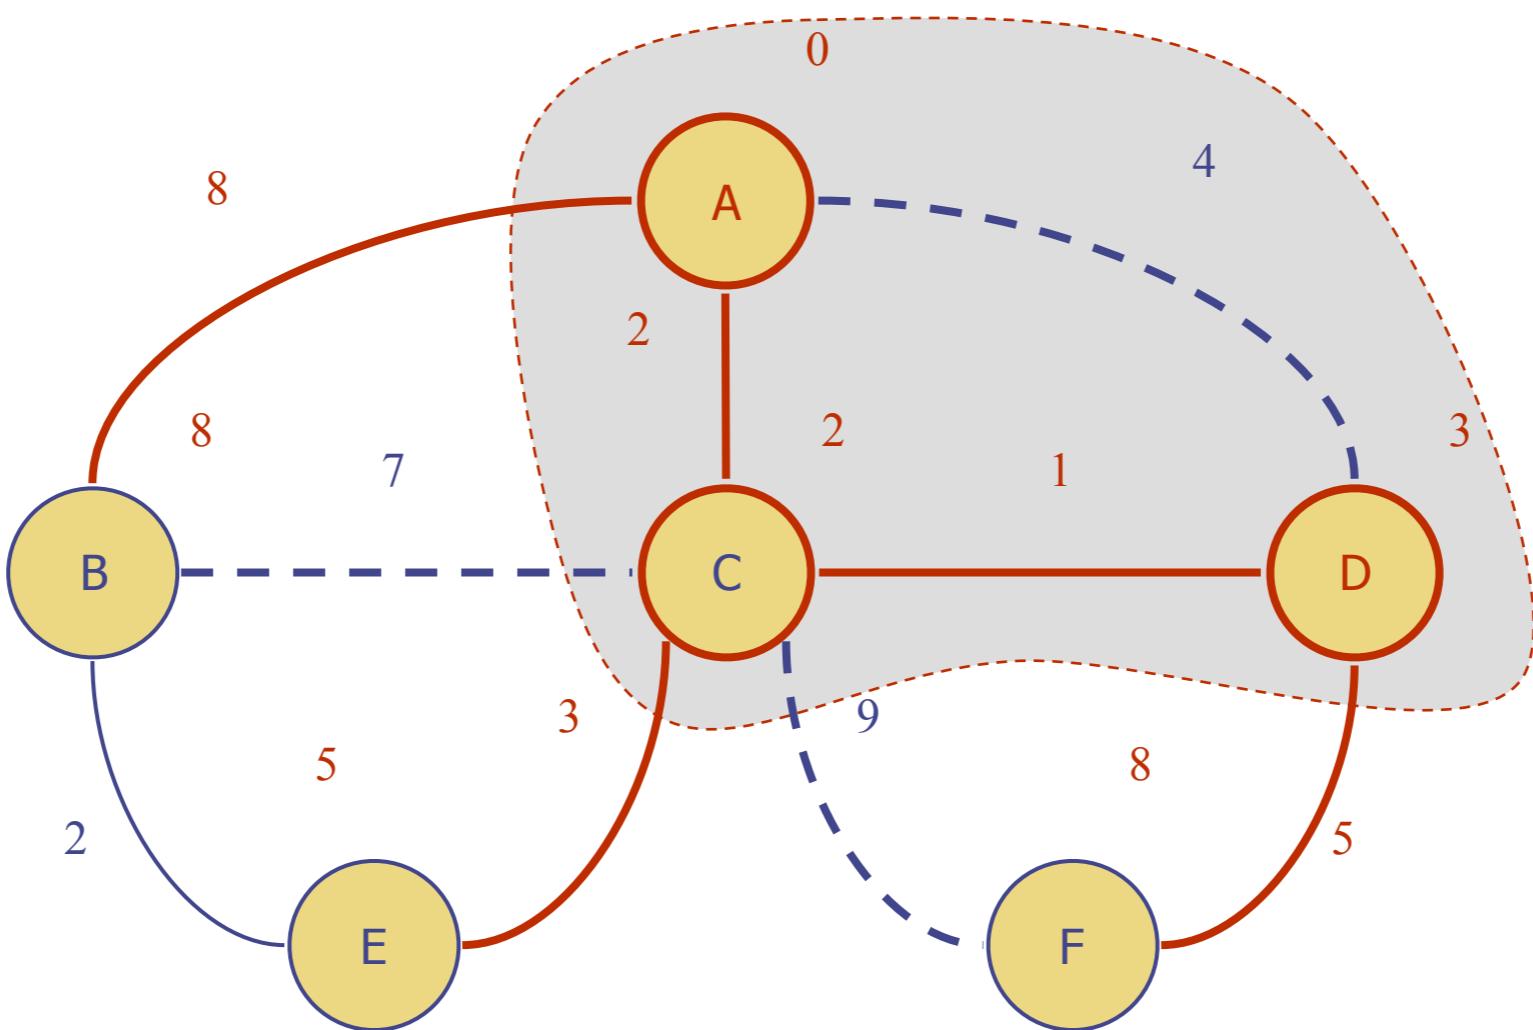
Example:

Tree of shortest paths from Providence



Providence airport

Shortest Paths



Dijkstra's Algorithm

- The distance of a vertex v from a vertex s is the length of a shortest path between s and v
- Dijkstra's algorithm computes the distances of all the vertices from a given start vertex s
- Assumptions:
 - the graph is connected
 - the edges are undirected
 - the edge weights are nonnegative

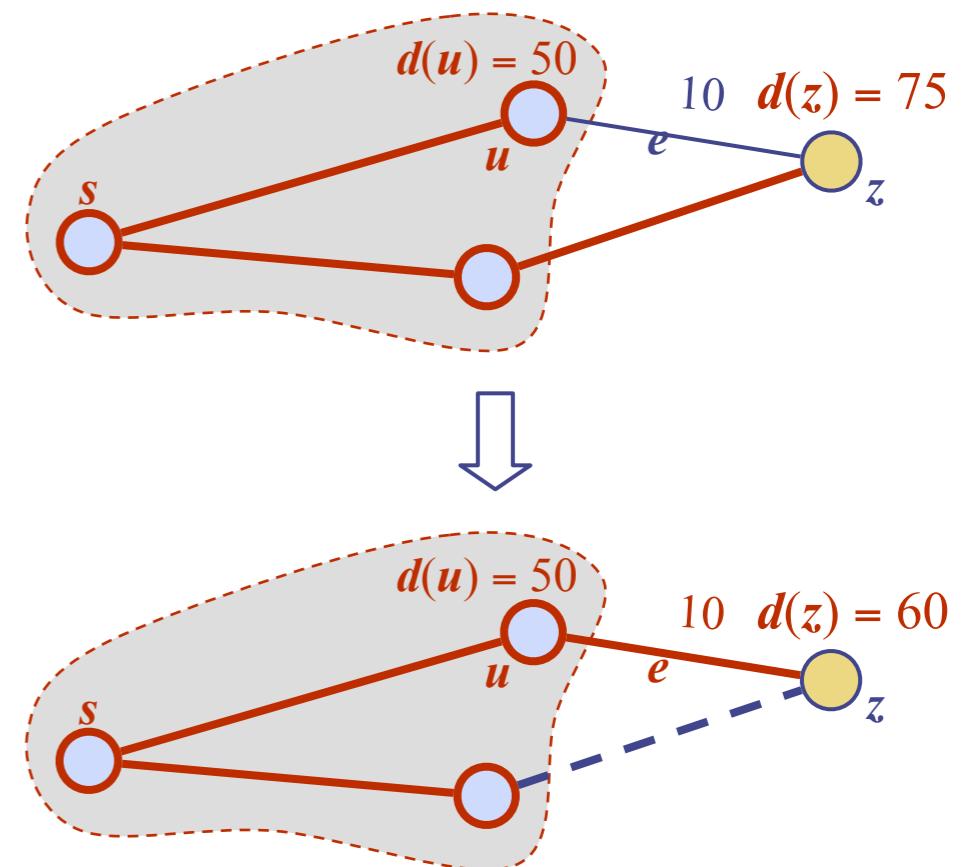
- We grow a “**cloud**” of vertices, beginning with s and eventually covering all the vertices
- We store with each vertex v a label $d(v)$ representing the distance of v from s in the subgraph consisting of the cloud and its adjacent vertices
- At each step
 - We add to the cloud the vertex u outside the cloud with the smallest distance label, $d(u)$
 - We update the labels of the vertices adjacent to u

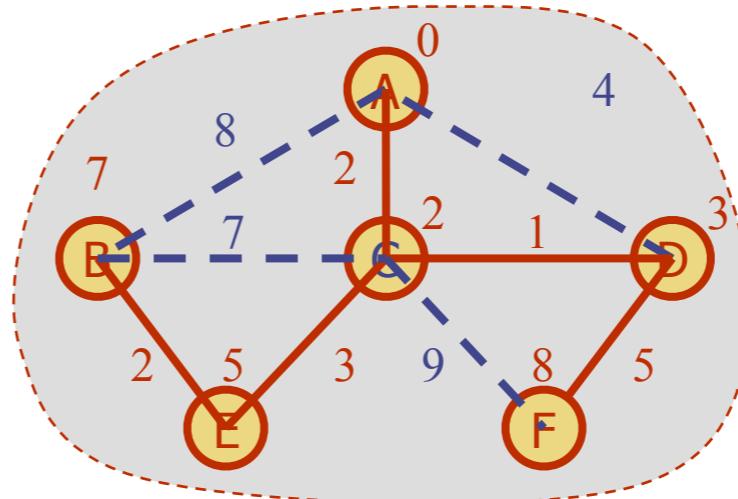
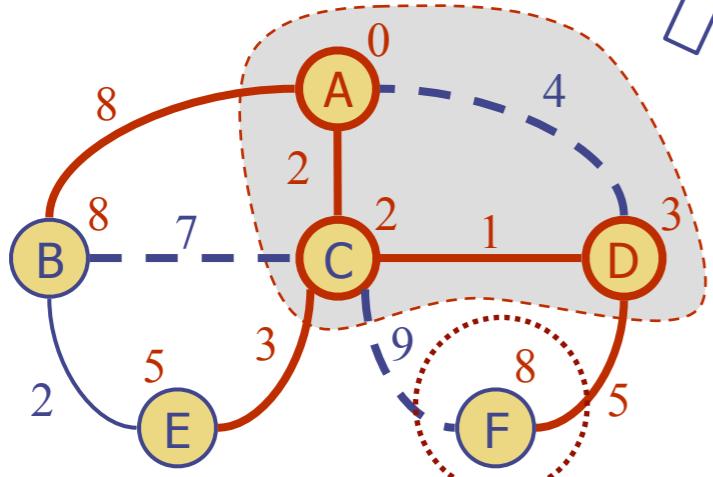
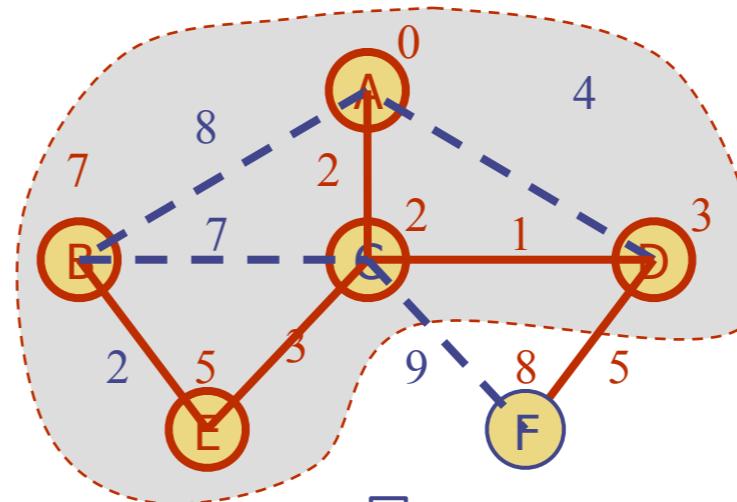
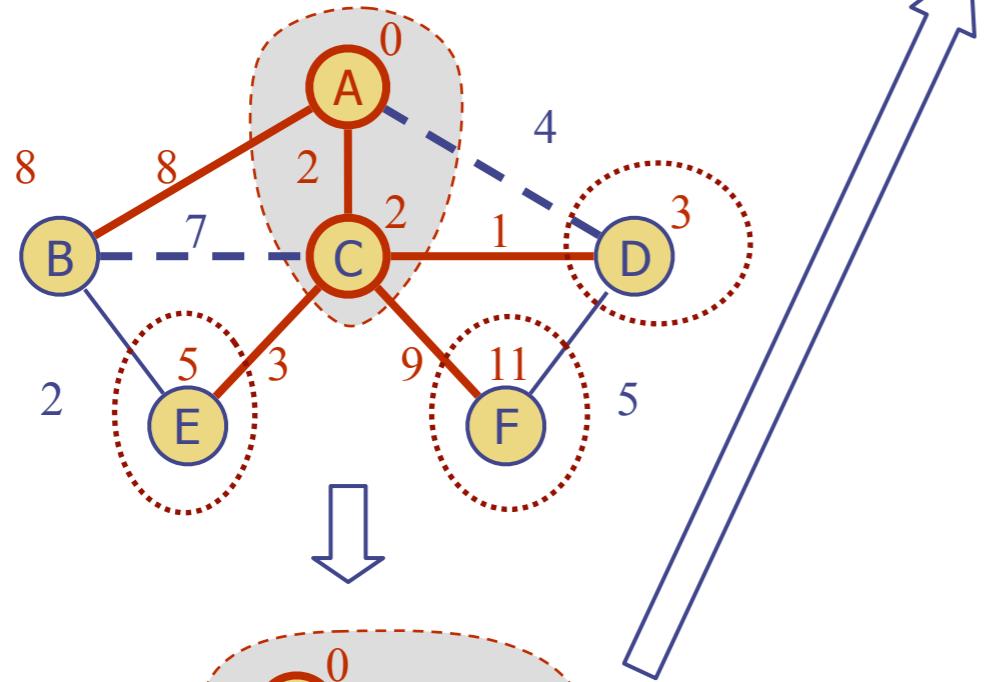
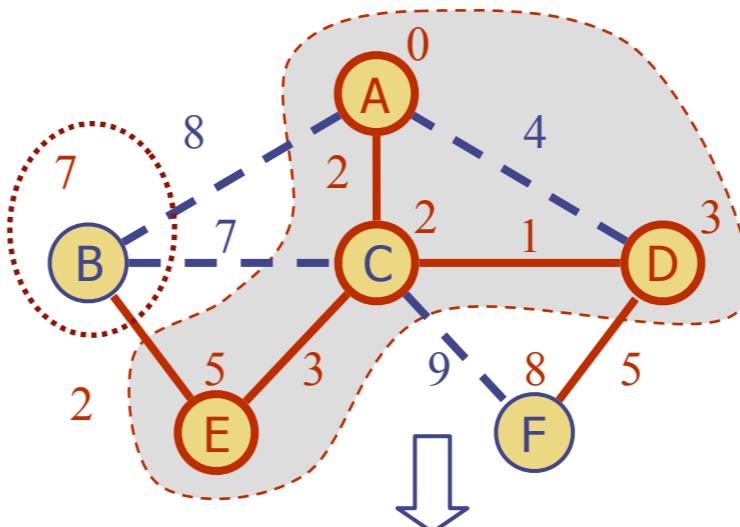
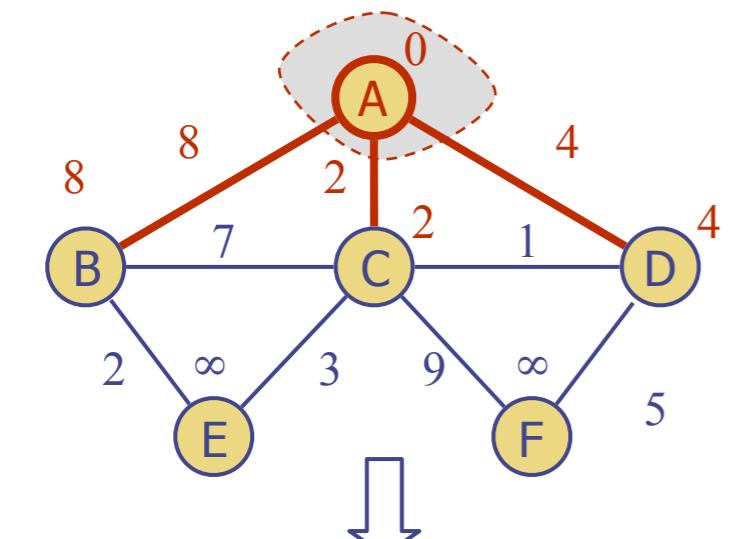
Dijkstra's Algorithm

- 1 Assign to every node a tentative distance value: set it to **zero** for our initial node and to **infinity** for all other nodes.
- 2 Set the initial node as current. Mark all other nodes unvisited. Create a set of all the unvisited nodes called the ***unvisited set***.
- 3 For the current node, **consider all of its unvisited neighbors and calculate their *tentative distances***. **Compare the newly calculated *tentative distance* to the current assigned value and assign the smaller one**. For example, if the current node *A* is marked with a distance of 6, and the edge connecting it with a neighbor *B* has length 2, then the distance to *B* (through *A*) will be $6 + 2 = 8$. If *B* was previously marked with a distance greater than 8 then change it to 8. Otherwise, keep the current value.
- 4 **When we are done considering all of the neighbors of the current node, mark the current node as visited and remove it from the *unvisited set*. A visited node will never be checked again.**
- 5 If the **destination node has been marked visited** (when planning a route between two specific nodes) **or if the smallest tentative distance among the nodes in the *unvisited set* is infinity** (when planning a complete traversal; occurs when there is no connection between the initial node and remaining unvisited nodes), **then stop**. The algorithm has finished.
- 6 Otherwise, **select the unvisited node that is marked with the smallest tentative distance, set it as the new "current node", and go back to step 3**.

Edge “Relaxation”

- Consider an edge $e = (u, z)$ such that
 - u is the vertex most recently added to the cloud
 - z is not in the cloud
- The “relaxation” of edge e updates distance $d(z)$ as follows:
$$d(z) \leftarrow \min\{d(z), d(u) + \text{weight}(e)\}$$





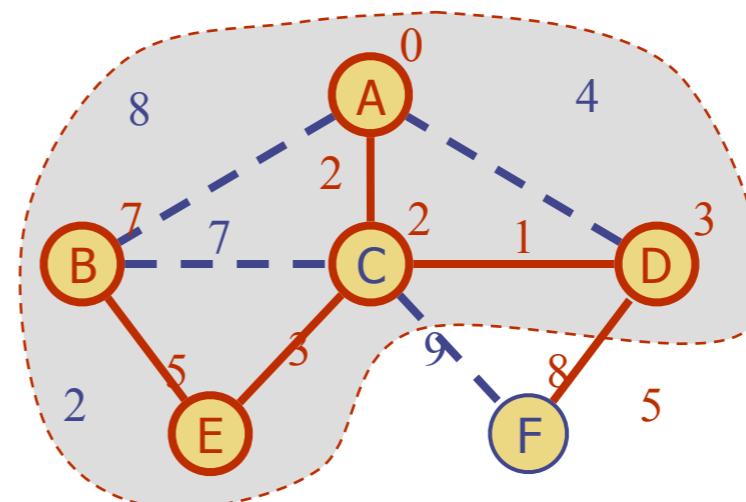
Analysis

- Graph operations
 - Method `incidentEdges` is called once for each vertex
- Label operations
 - We set/get the distance and locator labels of vertex z $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex in the priority queue is modified at most $\deg(w)$ times, where each key change takes $O(\log n)$ time
- Dijkstra's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time can also be expressed as $O(m \log n)$ since the graph is connected

Why Dijkstra's Algorithm Works

Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.

- Suppose it didn't find all shortest distances. Let F be the first wrong vertex the algorithm processed.
- When the previous node, D, on the true shortest path was considered, its distance was correct.
- But the edge (D,F) was **relaxed** at that time!
- Thus, so long as $d(F) \geq d(D)$, F's distance cannot be wrong. That is, there is no wrong vertex.



Dijkstra's algorithm: Correctness proof by induction

- G is the input graph,
- s is the source vertex,
- $\ell(u,v)$ is the length of an edge from u to v
- $\ell(Q)$ is the length of a path Q
- V is the set of vertices
- R is the set of visited nodes
- $d(v)$ is the current distance from the source, updated along the algorithm
- $\delta(v)$ be the shortest path distance from s -to- v (true distance)

DIJKSTRA(G, s)

for all $u \in V \setminus \{s\}$, $d(u) = \infty$

$d(s) = 0$

$R = \{\}$

while $R \neq V$

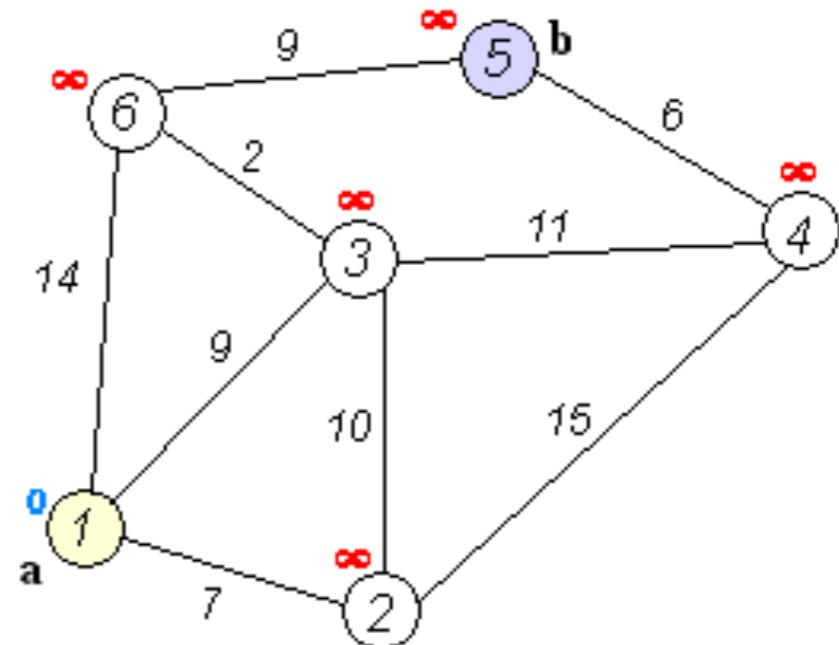
 pick $u \notin R$ with smallest $d(u)$

$R = R \cup \{u\}$

 for all vertices v adjacent to u

 if $d(v) > d(u) + \ell(u, v)$

$d(v) = d(u) + \ell(u, v)$



Let $d(v)$ be the label found by the algorithm and let $\delta(v)$ be the shortest path distance from s-to-v.

We want to show that $d(v) = \delta(v)$ for every vertex v at the end of the algorithm, showing that the algorithm correctly computes the distances.

We prove this by induction on the cardinality of R ($|R|$) (R is the set of visited nodes and its size (cardinality) increases by one at each step of the algorithm)

Lemma: For each $x \in R$, $d(x) = \delta(x)$.

Proof by Induction:

- 1) **Base case** ($|R| = 1$): Since R only grows in size, the only time $|R| = 1$ is when $R = \{s\}$ and $d(s) = 0 = \delta(s)$, which is correct.
- 2) **Inductive hypothesis (IH)** : Let u be the last vertex added to R . Let $R' = R \cup \{u\}$. Our I.H. is: for each $x \in R'$, $d(x) = \delta(x)$.

Idea: By the inductive hypothesis, for every vertex in R' that isn't u , we have the correct distance label. We need only show that $d(u) = \delta(u)$ to complete the proof (induction on the size of R)

We show that this is true by contradiction.

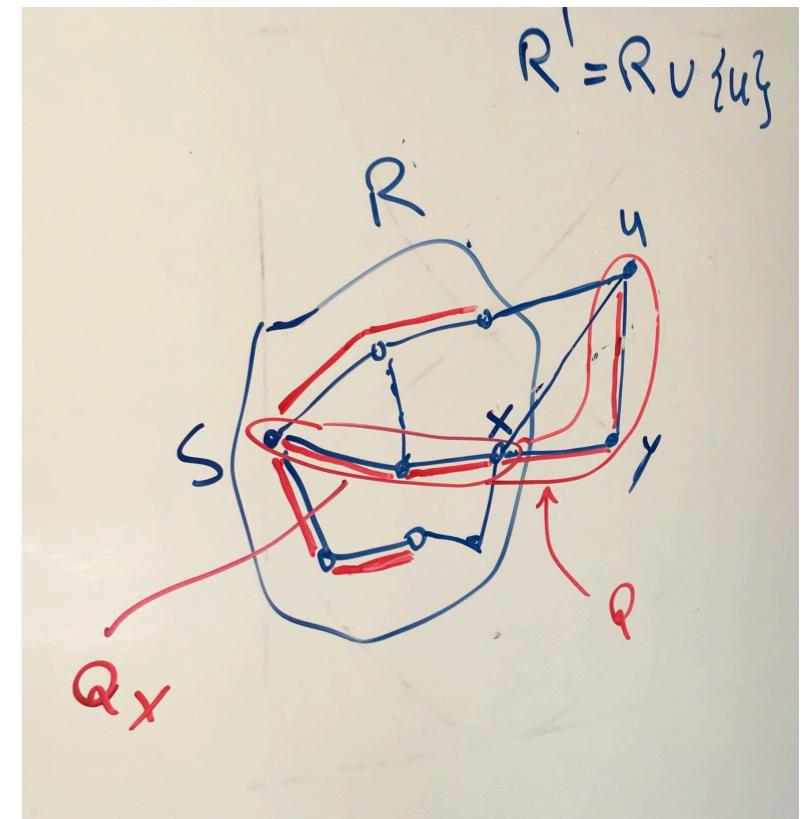
Suppose for a contradiction that the shortest path from s -to- u is Q and has length

$$\ell(Q) < d(u)$$

Q starts in R' and at some vertex leaves R' (to get to u). Let (x,y) be the first edge along Q that leaves R' .

Let Q_x be the s -to- x subpath of Q . Clearly:

$$\ell(Q_x) + \ell(x,y) \leq \ell(Q).$$



Since $\mathbf{d}(\mathbf{x})$ is the length of the shortest s-to-x path by the I.H., $\mathbf{d}(\mathbf{x}) = \ell(\mathbf{Q}_x)$, giving us

$$\mathbf{d}(\mathbf{x}) + \ell(\mathbf{x}, \mathbf{y}) \leq \ell(\mathbf{Q}).$$

Since \mathbf{y} is adjacent to \mathbf{x} , $\mathbf{d}(\mathbf{y})$ must have been updated by the algorithm, so

$$\mathbf{d}(\mathbf{y}) \leq \mathbf{d}(\mathbf{x}) + \ell(\mathbf{x}, \mathbf{y}).$$

Finally, since \mathbf{u} was picked by the algorithm, \mathbf{u} must have the smallest distance label:

$$\mathbf{d}(\mathbf{u}) \leq \mathbf{d}(\mathbf{y}).$$

Combining these inequalities in reverse order gives us the contradiction that $\mathbf{d}(\mathbf{x}) < \mathbf{d}(\mathbf{x})$.

Therefore, no such shorter path \mathbf{Q} can exist and so $\mathbf{d}(\mathbf{u}) = \delta(\mathbf{u})$.

This lemma shows the algorithm is correct by “applying” the lemma for $\mathbf{R} = \mathbf{V}$.

Steps to find the contradiction:

$$\mathbf{d(u) \leq d(y)}$$

using $d(y) \leq d(x) + \ell(x,y)$ we get

$$\mathbf{d(u) \leq d(x) + \ell(x,y)}$$

using $d(u) < \ell(Q)$ we get

$$\ell(Q) < d(x) + \ell(x,y)$$

using $\ell(Q_x) + \ell(x,y) \leq \ell(Q)$ we get

$$\ell(Q_x) + \ell(x,y) < d(x) + \ell(x,y)$$

$$\ell(Q_x) < d(x)$$

by the I.H. we have $\ell(Q_x) = d(x)$, hence

$$\mathbf{d(x) < d(x)}$$

contradiction!